


```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: train_data=pd.read_csv(r"C:\Users\manasa\Downloads\Mobile_Price_Classification
train_data
```

```
Out[2]:
```

	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	m_dep	mobile_wt
0	842	0	2.2	0	1	0	7	0.6	188
1	1021	1	0.5	1	0	1	53	0.7	136
2	563	1	0.5	1	2	1	41	0.9	145
3	615	1	2.5	0	0	0	10	0.8	131
4	1821	1	1.2	0	13	1	44	0.6	141
...
1995	794	1	0.5	1	0	1	2	0.8	106
1996	1965	1	2.6	1	0	0	39	0.2	187
1997	1911	0	0.9	1	1	1	36	0.7	108
1998	1512	0	0.9	0	4	1	46	0.1	145
1999	510	1	2.0	1	5	1	45	0.9	168

2000 rows × 21 columns



```
In [3]: test_data=pd.read_csv(r"C:\Users\manasa\Downloads\Mobile_Price_Classification_
test_data
```

```
Out[3]:
```

	id	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	m_dep	mobile
0	1	1043	1	1.8	1	14	0	5	0.1	
1	2	841	1	0.5	1	4	1	61	0.8	
2	3	1807	1	2.8	0	1	0	27	0.9	
3	4	1546	0	0.5	1	18	1	25	0.5	
4	5	1434	0	1.4	0	11	1	49	0.5	
...
995	996	1700	1	1.9	0	0	1	54	0.5	
996	997	609	0	1.8	1	0	0	13	0.9	
997	998	1185	0	1.4	0	1	1	8	0.5	
998	999	1533	1	0.5	1	0	0	50	0.4	
999	1000	1270	1	0.5	0	4	1	35	0.1	

1000 rows × 21 columns



```
In [4]: train_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2000 entries, 0 to 1999
Data columns (total 21 columns):
#   Column          Non-Null Count  Dtype
---  -
0   battery_power    2000 non-null   int64
1   blue             2000 non-null   int64
2   clock_speed      2000 non-null   float64
3   dual_sim         2000 non-null   int64
4   fc               2000 non-null   int64
5   four_g           2000 non-null   int64
6   int_memory       2000 non-null   int64
7   m_dep            2000 non-null   float64
8   mobile_wt        2000 non-null   int64
9   n_cores          2000 non-null   int64
10  pc               2000 non-null   int64
11  px_height        2000 non-null   int64
12  px_width         2000 non-null   int64
13  ram              2000 non-null   int64
14  sc_h             2000 non-null   int64
15  sc_w             2000 non-null   int64
16  talk_time        2000 non-null   int64
17  three_g          2000 non-null   int64
18  touch_screen     2000 non-null   int64
19  wifi             2000 non-null   int64
20  price_range      2000 non-null   int64
dtypes: float64(2), int64(19)
memory usage: 328.3 KB
```

```
In [5]: test_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 21 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   id              1000 non-null   int64
 1   battery_power   1000 non-null   int64
 2   blue            1000 non-null   int64
 3   clock_speed     1000 non-null   float64
 4   dual_sim        1000 non-null   int64
 5   fc              1000 non-null   int64
 6   four_g         1000 non-null   int64
 7   int_memory      1000 non-null   int64
 8   m_dep           1000 non-null   float64
 9   mobile_wt       1000 non-null   int64
10   n_cores         1000 non-null   int64
11   pc              1000 non-null   int64
12   px_height       1000 non-null   int64
13   px_width        1000 non-null   int64
14   ram             1000 non-null   int64
15   sc_h            1000 non-null   int64
16   sc_w            1000 non-null   int64
17   talk_time       1000 non-null   int64
18   three_g         1000 non-null   int64
19   touch_screen    1000 non-null   int64
20   wifi            1000 non-null   int64
dtypes: float64(2), int64(19)
memory usage: 164.2 KB
```

```
In [6]: x=train_data.drop('wifi',axis=1)
        y=train_data['wifi']
```

```
In [7]: x=test_data.drop('wifi',axis=1)
        y=test_data['wifi']
```

```
In [8]: train_data['dual_sim'].value_counts()
```

```
Out[8]: dual_sim
1      1019
0       981
Name: count, dtype: int64
```

```
In [9]: test_data['dual_sim'].value_counts()
```

```
Out[9]: dual_sim
1       517
0       483
Name: count, dtype: int64
```

```
In [10]: TG={"three_g":{"yes":1,"No":0}}
train_data=train_data.replace(TG)
print(train_data)
```

	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory
0	842	0	2.2	0	1	0	7
1	1021	1	0.5	1	0	1	53
2	563	1	0.5	1	2	1	41
3	615	1	2.5	0	0	0	10
4	1821	1	1.2	0	13	1	44
...
1995	794	1	0.5	1	0	1	2
1996	1965	1	2.6	1	0	0	39
1997	1911	0	0.9	1	1	1	36
1998	1512	0	0.9	0	4	1	46
1999	510	1	2.0	1	5	1	45

	m_dep	mobile_wt	n_cores	...	px_height	px_width	ram	sc_h	sc_w
0	0.6	188	2	...	20	756	2549	9	7
1	0.7	136	3	...	905	1988	2631	17	3
2	0.9	145	5	...	1263	1716	2603	11	2
3	0.8	131	6	...	1216	1786	2769	16	8
4	0.6	141	2	...	1208	1212	1411	8	2
...
1995	0.8	106	6	...	1222	1890	668	13	4
1996	0.2	187	4	...	915	1965	2032	11	10
1997	0.7	108	8	...	868	1632	3057	9	1
1998	0.1	145	5	...	336	670	869	18	10
1999	0.9	168	6	...	483	754	3919	19	4

	talk_time	three_g	touch_screen	wifi	price_range
0	19	0	0	1	1
1	7	1	1	0	2
2	9	1	1	0	2
3	11	1	0	0	2
4	15	1	1	0	1
...
1995	19	1	1	0	0
1996	16	1	1	1	2
1997	5	1	1	0	3
1998	19	1	1	1	0
1999	2	1	1	1	3

[2000 rows x 21 columns]

```
In [11]: TG={"three_g":{"yes":1,"No":0}}
test_data=test_data.replace(TG)
print(test_data)
```

	id	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory
0	1	1043	1	1.8	1	14	0	5
\								
1	2	841	1	0.5	1	4	1	61
2	3	1807	1	2.8	0	1	0	27
3	4	1546	0	0.5	1	18	1	25
4	5	1434	0	1.4	0	11	1	49
..
995	996	1700	1	1.9	0	0	1	54
996	997	609	0	1.8	1	0	0	13
997	998	1185	0	1.4	0	1	1	8
998	999	1533	1	0.5	1	0	0	50
999	1000	1270	1	0.5	0	4	1	35

	m_dep	mobile_wt	...	pc	px_height	px_width	ram	sc_h	sc_w	\
0	0.1	193	...	16	226	1412	3476	12	7	
1	0.8	191	...	12	746	857	3895	6	0	
2	0.9	186	...	4	1270	1366	2396	17	10	
3	0.5	96	...	20	295	1752	3893	10	0	
4	0.5	108	...	18	749	810	1773	15	8	
..	
995	0.5	170	...	17	644	913	2121	14	8	
996	0.9	186	...	2	1152	1632	1933	8	1	
997	0.5	80	...	12	477	825	1223	5	0	
998	0.4	171	...	12	38	832	2509	15	11	
999	0.1	140	...	19	457	608	2828	9	2	

	talk_time	three_g	touch_screen	wifi
0	2	0	1	0
1	7	1	0	0
2	10	0	1	1
3	7	1	1	0
4	7	1	0	1
..
995	15	1	1	0
996	19	0	1	1
997	14	1	0	0
998	6	0	1	0
999	3	1	0	1

[1000 rows x 21 columns]

```
In [12]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,train_size=0.7,random_state
x_train.shape,x_test.shape)
```

Out[12]: ((700, 20), (300, 20))

```
In [13]: from sklearn.ensemble import RandomForestClassifier  
rfc=RandomForestClassifier()  
rfc.fit(x_train,y_train)
```

Out[13]: RandomForestClassifier()

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [15]: rf=RandomForestClassifier()  
params={'max_depth':[2,3,5,10,20],  
'min_samples_leaf':[5,10,20.50,100,200],  
'n_estimators':[10,25,30,50,100,200]}
```

```
In [21]: from sklearn.model_selection import GridSearchCV  
grid_search=GridSearchCV(estimator=rf,param_grid=params,cv=2,scoring='accuracy')  
grid_search.fit(x_train,y_train)
```

C:\Users\manasa\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\model_selection_validation.py:378: FitFailedWarning:
60 fits failed out of a total of 300.
The score on these train-test partitions for these parameters will be set to nan.
If these failures are not expected, you can try to debug them by setting error_score='raise'.

Below are more details about the failures:

```
-----
---
60 fits failed with the following error:
Traceback (most recent call last):
  File "C:\Users\manasa\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\model_selection\_validation.py", line 686, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "C:\Users\manasa\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\ensemble\_forest.py", line 340, in fit
    self._validate_params()
  File "C:\Users\manasa\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\base.py", line 600, in _validate_params
    validate_parameter_constraints(
  File "C:\Users\manasa\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\utils\_param_validation.py", line 97, in validate_parameter_constraints
    raise InvalidParameterError(
sklearn.utils._param_validation.InvalidParameterError: The 'min_samples_leaf' parameter of RandomForestClassifier must be an int in the range [1, inf) or a float in the range (0.0, 1.0). Got 20.5 instead.

warnings.warn(some_fits_failed_message, FitFailedWarning)
C:\Users\manasa\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\model_selection\_search.py:952: UserWarning: One or more of the test scores are non-finite: [0.50714286 0.49428571 0.52      0.51285714 0.49857143
0.48857143
0.51285714 0.48857143 0.48857143 0.49714286 0.50571429 0.49285714
nan nan nan nan nan nan
0.48428571 0.50285714 0.49285714 0.51714286 0.5      0.5
0.5      0.5      0.5      0.5      0.5      0.5
0.48285714 0.49571429 0.51285714 0.5      0.49428571 0.51
0.47857143 0.52142857 0.47285714 0.5      0.49428571 0.48714286
nan nan nan nan nan nan
0.48571429 0.51285714 0.48428571 0.49      0.49857143 0.50142857
0.5      0.5      0.5      0.5      0.5      0.5
0.48428571 0.51428571 0.49428571 0.49714286 0.51714286 0.48857143
0.50714286 0.48428571 0.49571429 0.49857143 0.49857143 0.49857143
nan nan nan nan nan nan
0.48857143 0.48142857 0.50428571 0.50857143 0.49714286 0.50142857
0.5      0.5      0.5      0.5      0.5      0.5
0.50428571 0.47571429 0.51285714 0.48714286 0.46571429 0.51285714
0.52285714 0.51      0.50857143 0.49142857 0.49571429 0.50428571
nan nan nan nan nan nan
0.50428571 0.48285714 0.48285714 0.48571429 0.50857143 0.5
0.5      0.5      0.5      0.5      0.5      0.5
0.51857143 0.50428571 0.49428571 0.49285714 0.52428571 0.49142857
0.52      0.50285714 0.50714286 0.50714286 0.48714286 0.50428571
nan nan nan nan nan nan
```



```
0.45714286 0.50428571 0.48714286 0.48857143 0.49285714 0.49571429
0.5         0.5         0.5         0.5         0.5         0.5         ]
warnings.warn(
```

```
Out[21]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
                    param_grid={'max_depth': [2, 3, 5, 10, 20],
                                'min_samples_leaf': [5, 10, 20.5, 100, 200],
                                'n_estimators': [10, 25, 30, 50, 100, 200]},
                    scoring='accuracy')
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [22]: grid_search.best_score_
```

```
Out[22]: 0.5242857142857142
```

```
In [23]: grid_search.fit(x_train,y_train)
```

C:\Users\manasa\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\model_selection_validation.py:378: FitFailedWarning:
60 fits failed out of a total of 300.
The score on these train-test partitions for these parameters will be set to nan.
If these failures are not expected, you can try to debug them by setting error_score='raise'.

Below are more details about the failures:

```
-----
---
60 fits failed with the following error:
Traceback (most recent call last):
  File "C:\Users\manasa\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\model_selection\_validation.py", line 686, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "C:\Users\manasa\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\ensemble\_forest.py", line 340, in fit
    self._validate_params()
  File "C:\Users\manasa\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\base.py", line 600, in _validate_params
    validate_parameter_constraints(
  File "C:\Users\manasa\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\utils\_param_validation.py", line 97, in validate_parameter_constraints
    raise InvalidParameterError(
sklearn.utils._param_validation.InvalidParameterError: The 'min_samples_leaf' parameter of RandomForestClassifier must be an int in the range [1, inf) or a float in the range (0.0, 1.0). Got 20.5 instead.
```

```
warnings.warn(some_fits_failed_message, FitFailedWarning)
C:\Users\manasa\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\model_selection\_search.py:952: UserWarning: One or more of the test scores are non-finite: [0.49857143 0.47142857 0.51857143 0.51      0.49857143
0.50571429
0.48857143 0.5          0.51285714 0.50142857 0.49857143 0.49428571
nan          nan          nan          nan          nan          nan
0.51428571 0.50714286 0.48428571 0.49          0.50285714 0.5
0.5          0.5          0.5          0.5          0.5          0.5
0.49857143 0.53142857 0.47285714 0.49          0.51142857 0.48142857
0.54          0.49285714 0.51285714 0.50285714 0.47          0.49
nan          nan          nan          nan          nan          nan
0.45714286 0.48857143 0.50857143 0.50285714 0.49857143 0.50714286
0.5          0.5          0.5          0.5          0.5          0.5
0.50285714 0.5          0.53714286 0.50142857 0.51142857 0.49142857
0.50428571 0.50285714 0.51          0.49857143 0.47857143 0.49714286
nan          nan          nan          nan          nan          nan
0.49571429 0.46714286 0.50285714 0.47857143 0.48285714 0.49571429
0.5          0.5          0.5          0.5          0.5          0.5
0.49          0.50571429 0.5          0.50285714 0.52142857 0.50142857
0.51142857 0.49714286 0.48428571 0.48857143 0.49428571 0.48142857
nan          nan          nan          nan          nan          nan
0.46714286 0.46857143 0.5          0.49857143 0.48571429 0.48428571
0.5          0.5          0.5          0.5          0.5          0.5
0.49714286 0.50142857 0.5          0.49571429 0.48714286 0.50285714
0.51285714 0.46          0.49571429 0.48857143 0.49142857 0.50714286
nan          nan          nan          nan          nan          nan
```

```
0.51428571 0.47571429 0.48857143 0.49142857 0.48428571 0.48571429
0.5         0.5         0.5         0.5         0.5         0.5         ]
warnings.warn(
```

```
Out[23]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
                    param_grid={'max_depth': [2, 3, 5, 10, 20],
                                'min_samples_leaf': [5, 10, 20.5, 100, 200],
                                'n_estimators': [10, 25, 30, 50, 100, 200]},
                    scoring='accuracy')
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [24]: grid_search.best_score_
```

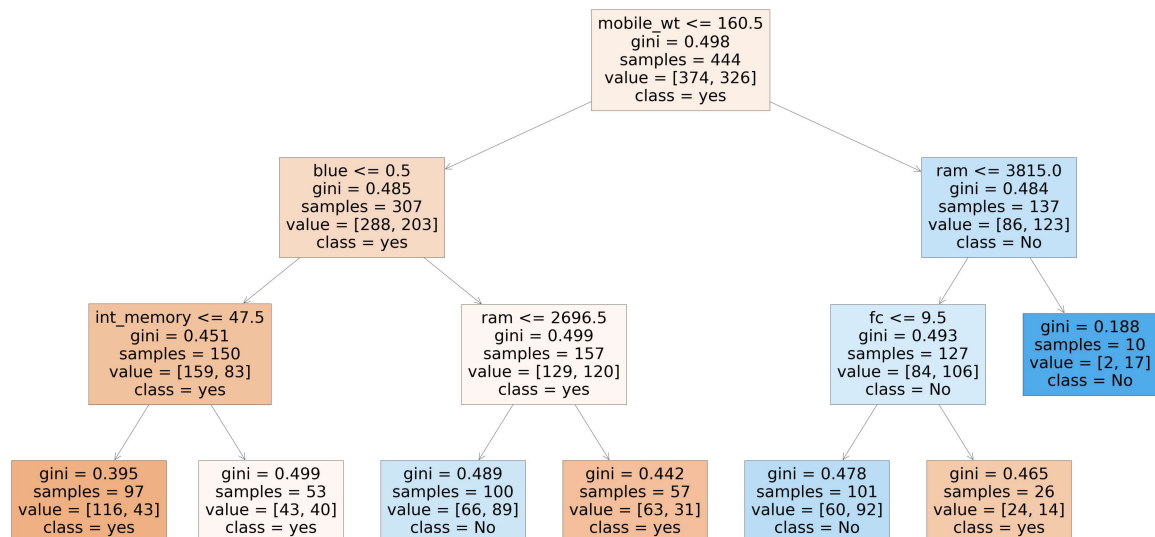
```
Out[24]: 0.54
```

```
In [25]: rf_best=grid_search.best_estimator_
         print(rf_best)
```

```
RandomForestClassifier(max_depth=3, min_samples_leaf=10, n_estimators=10)
```

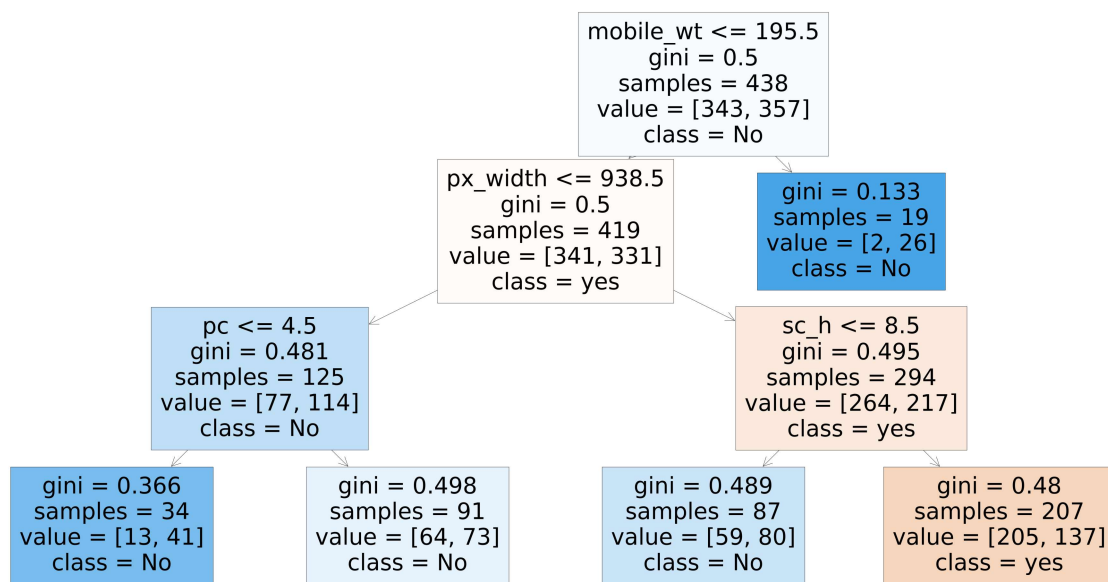
```
In [27]: from sklearn.tree import plot_tree
plt.figure(figsize=(80,40))
plot_tree(rf_best.estimators_[5],feature_names=x.columns,class_names=["yes","No"],
```

```
Out[27]: [Text(0.5769230769230769, 0.875, 'mobile_wt <= 160.5\ngini = 0.498\nsamples = 444\nvalue = [374, 326]\nnclass = yes'),
Text(0.3076923076923077, 0.625, 'blue <= 0.5\ngini = 0.485\nsamples = 307\nvalue = [288, 203]\nnclass = yes'),
Text(0.15384615384615385, 0.375, 'int_memory <= 47.5\ngini = 0.451\nsamples = 150\nvalue = [159, 83]\nnclass = yes'),
Text(0.07692307692307693, 0.125, 'gini = 0.395\nsamples = 97\nvalue = [116, 43]\nnclass = yes'),
Text(0.23076923076923078, 0.125, 'gini = 0.499\nsamples = 53\nvalue = [43, 40]\nnclass = yes'),
Text(0.46153846153846156, 0.375, 'ram <= 2696.5\ngini = 0.499\nsamples = 157\nvalue = [129, 120]\nnclass = yes'),
Text(0.38461538461538464, 0.125, 'gini = 0.489\nsamples = 100\nvalue = [66, 89]\nnclass = No'),
Text(0.5384615384615384, 0.125, 'gini = 0.442\nsamples = 57\nvalue = [63, 31]\nnclass = yes'),
Text(0.8461538461538461, 0.625, 'ram <= 3815.0\ngini = 0.484\nsamples = 137\nvalue = [86, 123]\nnclass = No'),
Text(0.7692307692307693, 0.375, 'fc <= 9.5\ngini = 0.493\nsamples = 127\nvalue = [84, 106]\nnclass = No'),
Text(0.6923076923076923, 0.125, 'gini = 0.478\nsamples = 101\nvalue = [60, 92]\nnclass = No'),
Text(0.8461538461538461, 0.125, 'gini = 0.465\nsamples = 26\nvalue = [24, 14]\nnclass = yes'),
Text(0.9230769230769231, 0.375, 'gini = 0.188\nsamples = 10\nvalue = [2, 17]\nnclass = No')]
```



```
In [28]: from sklearn.tree import plot_tree
plt.figure(figsize=(80,40))
plot_tree(rf_best.estimators_[7],feature_names=x.columns,class_names=["yes","No"])
```

```
Out[28]: [Text(0.625, 0.875, 'mobile_wt <= 195.5\nngini = 0.5\nsamples = 438\nvalue = [343, 357]\nnclass = No'),
Text(0.5, 0.625, 'px_width <= 938.5\nngini = 0.5\nsamples = 419\nvalue = [341, 331]\nnclass = yes'),
Text(0.25, 0.375, 'pc <= 4.5\nngini = 0.481\nsamples = 125\nvalue = [77, 114]\nnclass = No'),
Text(0.125, 0.125, 'gini = 0.366\nsamples = 34\nvalue = [13, 41]\nnclass = No'),
Text(0.375, 0.125, 'gini = 0.498\nsamples = 91\nvalue = [64, 73]\nnclass = No'),
Text(0.75, 0.375, 'sc_h <= 8.5\nngini = 0.495\nsamples = 294\nvalue = [264, 217]\nnclass = yes'),
Text(0.625, 0.125, 'gini = 0.489\nsamples = 87\nvalue = [59, 80]\nnclass = No'),
Text(0.875, 0.125, 'gini = 0.48\nsamples = 207\nvalue = [205, 137]\nnclass = yes'),
Text(0.75, 0.625, 'gini = 0.133\nsamples = 19\nvalue = [2, 26]\nnclass = No')]
```



```
In [29]: rf_best.feature_importances_
```

```
Out[29]: array([0.04650189, 0.10299328, 0.01439367, 0.005532626, 0.01964821, 0.13808652, 0.03685331, 0.09960983,
0.03901548, 0.09029733, 0.04591778, 0.10955797, 0.10997793,
0.0477564 , 0.0169304 , 0.00460595, 0.02252778])
```

```
In [35]: imp_df=pd.DataFrame({"varname":x_train.columns,"Imp":rf_best.feature_importance_})
imp_df.sort_values(by="Imp",ascending=False)
```

Out[35]:

	varname	Imp
7	int_memory	0.138087
14	ram	0.109978
13	px_width	0.109558
1	battery_power	0.102993
9	mobile_wt	0.099610
11	pc	0.090297
5	fc	0.055326
15	sc_h	0.047756
0	id	0.046502
12	px_height	0.045918
10	n_cores	0.039015
8	m_dep	0.036853
19	touch_screen	0.022528
6	four_g	0.019648
17	talk_time	0.016930
2	blue	0.014394
18	three_g	0.004606
4	dual_sim	0.000000
3	clock_speed	0.000000
16	sc_w	0.000000

In []: